

The Z Garbage Collector

Low Latency GC for OpenJDK



Per Lidén & Stefan Karlsson
HotSpot Garbage Collection Team
Jfokus VM Tech Summit 2018

Safe Harbor Statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.

Agenda

- 1 ➤ What is ZGC?
- 2 ➤ Some Numbers
- 3 ➤ Under The Hood
- 4 ➤ Going Forward
- 5 ➤ How To Get Started

Agenda

- 1 What is ZGC?
- 2 Some Numbers
- 3 Under The Hood
- 4 Going Forward
- 5 How To Get Started

A Scalable Low Latency Garbage Collector

Goals

TB

Multi-terabyte heaps

10_{ms}

Max GC pause time



Lay the foundation for
future GC features

15%

Max application
throughput reduction

GC pause times **do not** increase with heap or live-set size

At a Glance

- New garbage collector
- Load barriers
- Colored pointers
- Single generation
- Partial compaction
- Region-based
- Immediate memory reuse
- NUMA-aware

- Concurrent
 - ✓ Marking
 - ✓ Relocation/Compaction
 - ✓ Relocation Set Selection
 - ✓ Reference Processing
 - ✓ JNI WeakRefs Cleaning
 - StringTable/SymbolTable Cleaning
 - Class Unloading



Current Status

- Design and implementation approaching mature and stable
- Main focus on **Linux/x86_64**
 - Other platforms can be added if there's enough demand
- Performance looks very good
 - Both in terms of latency and throughput




Agenda

- 1 What is ZGC?
- 2 Some Numbers**
- 3 Under The Hood
- 4 Going Forward
- 5 How To Get Started

SPECjbb[®]2015 – Score

(Higher is better)



Mode: Composite

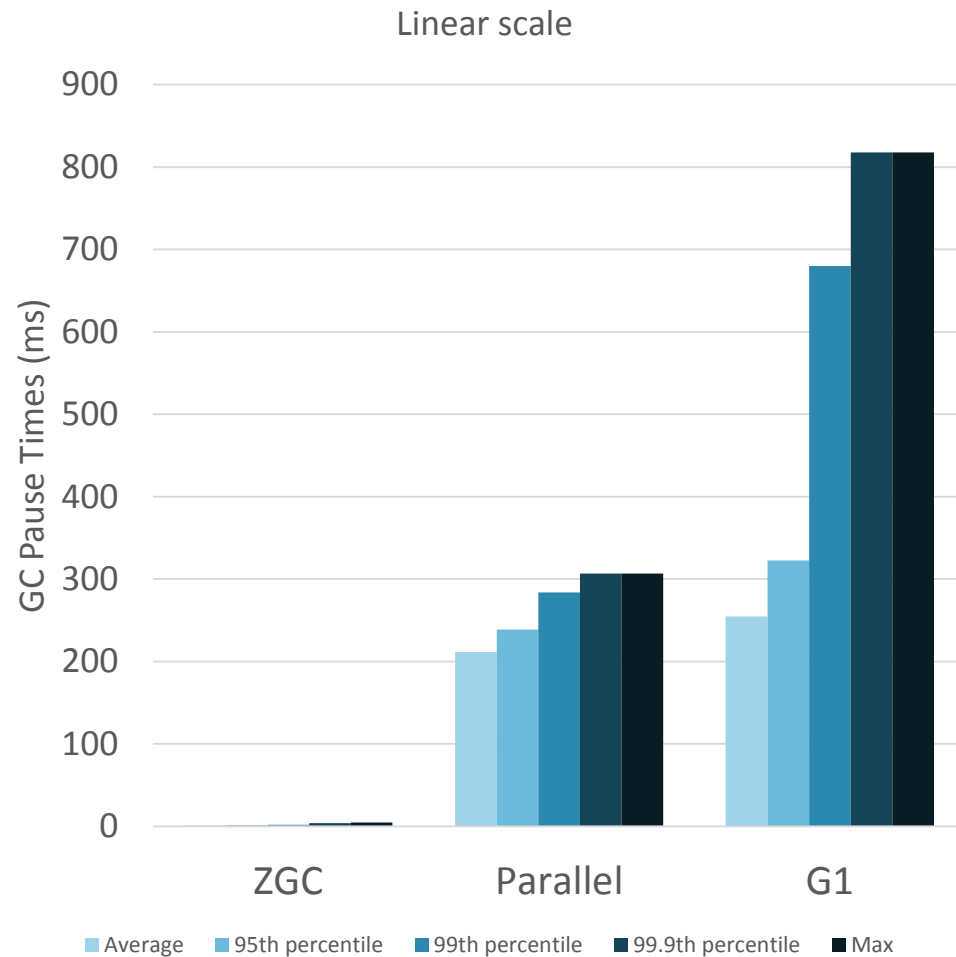
Heap Size: 128G

OS: Oracle Linux 7.4



HW: Intel Xeon E5-2690 2.9GHz
2 sockets, 16 cores (32 hw-threads)

SPECjbb[®]2015 is a registered trademark of the Standard Performance Evaluation Corporation (spec.org). The actual results are not represented as compliant because the SUT may not meet SPEC's requirements for general availability.

SPECjbb® 2015 – Pause Times



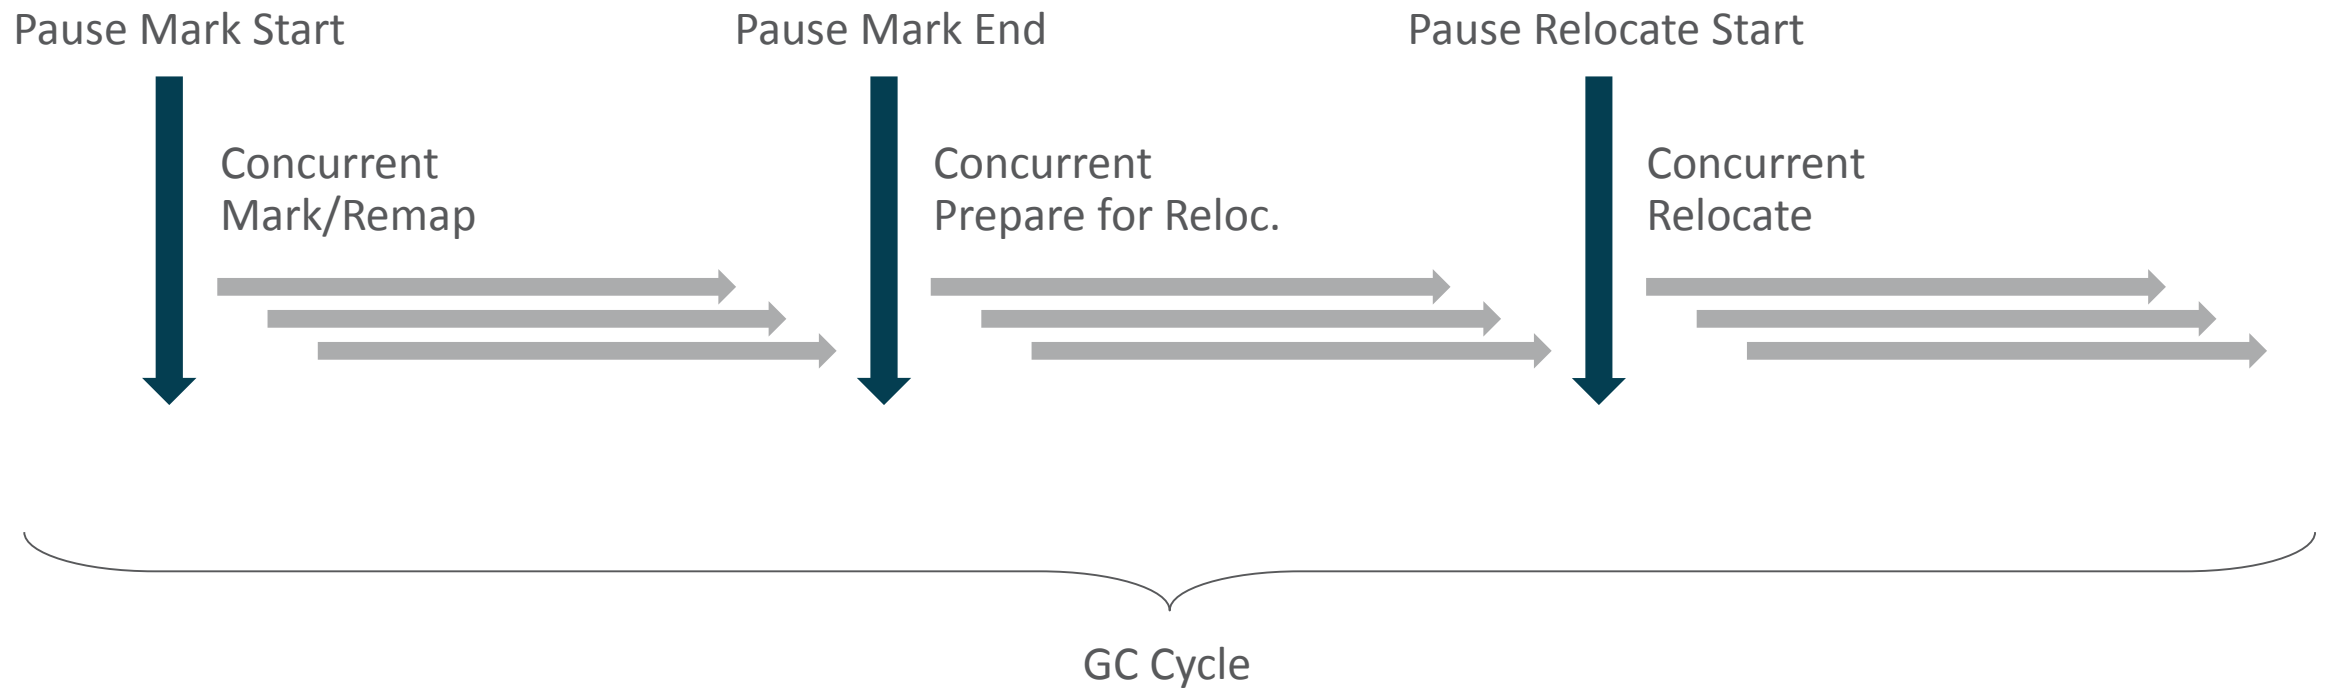
(Lower is better)



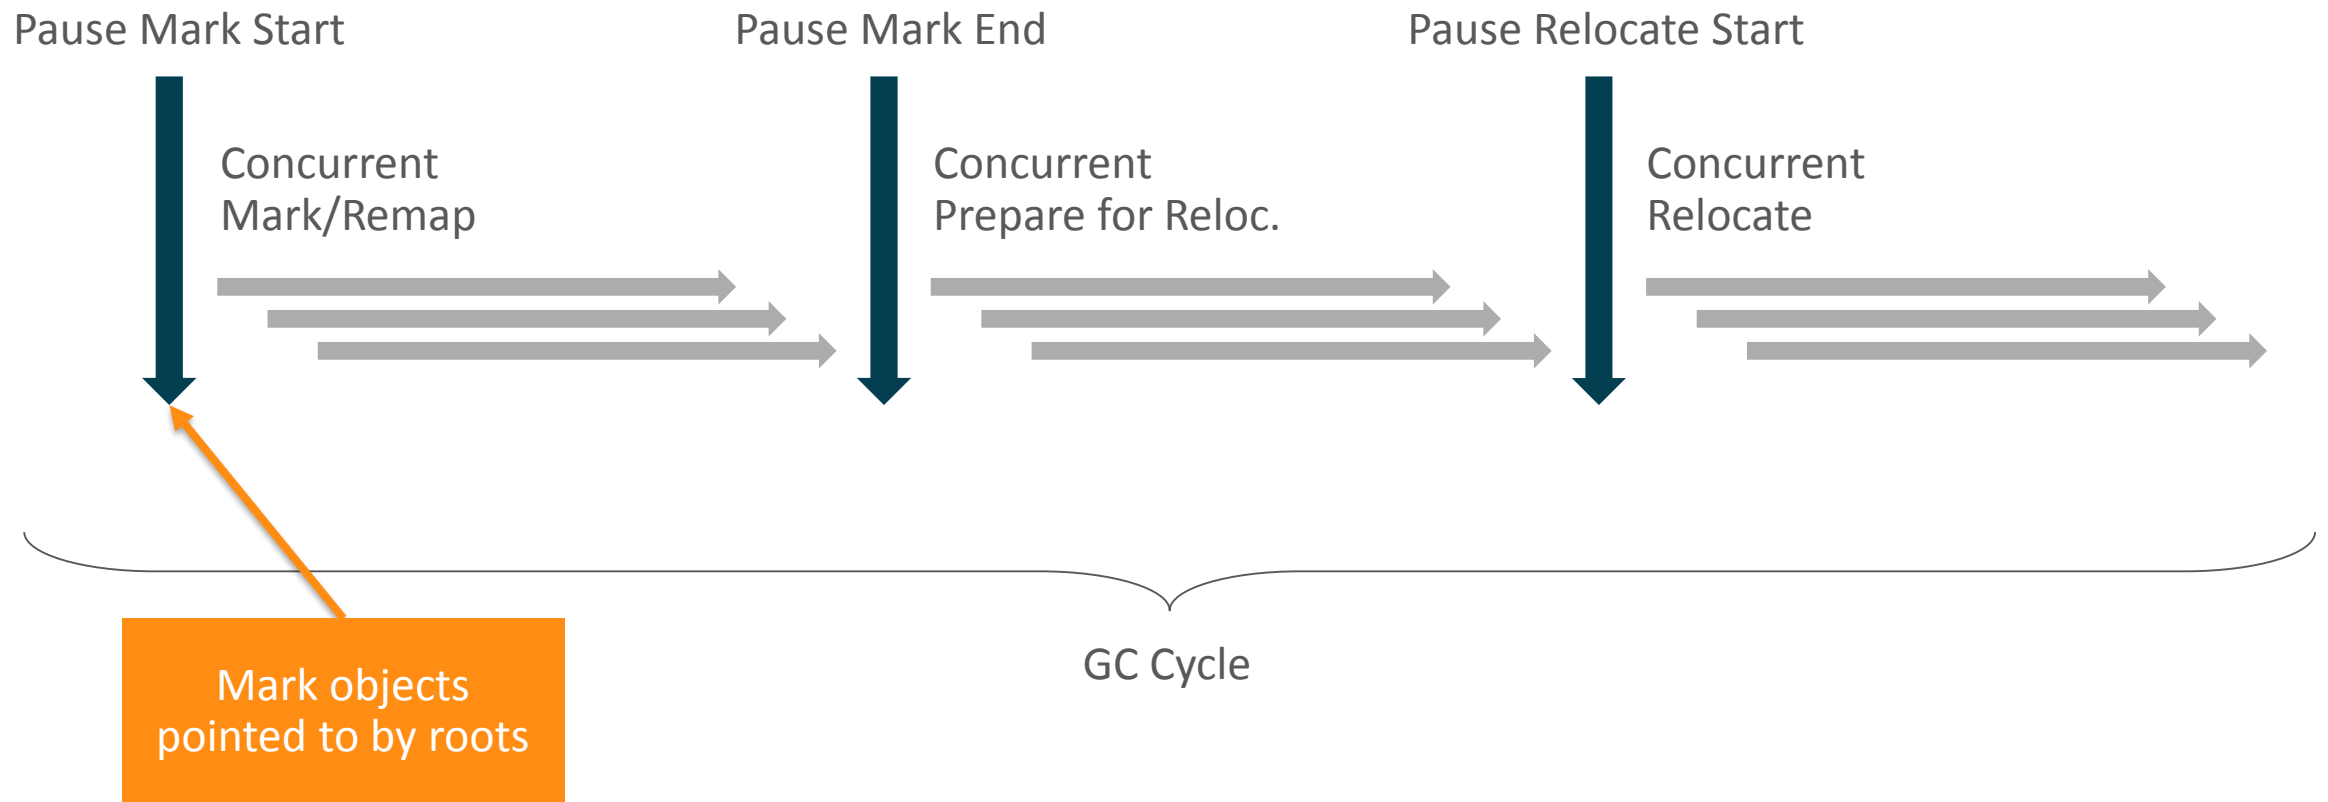
Agenda

- 1 ➤ What is ZGC?
- 2 ➤ Some Numbers
- 3 ➤ Under The Hood**
- 4 ➤ Going Forward
- 5 ➤ How To Get Started

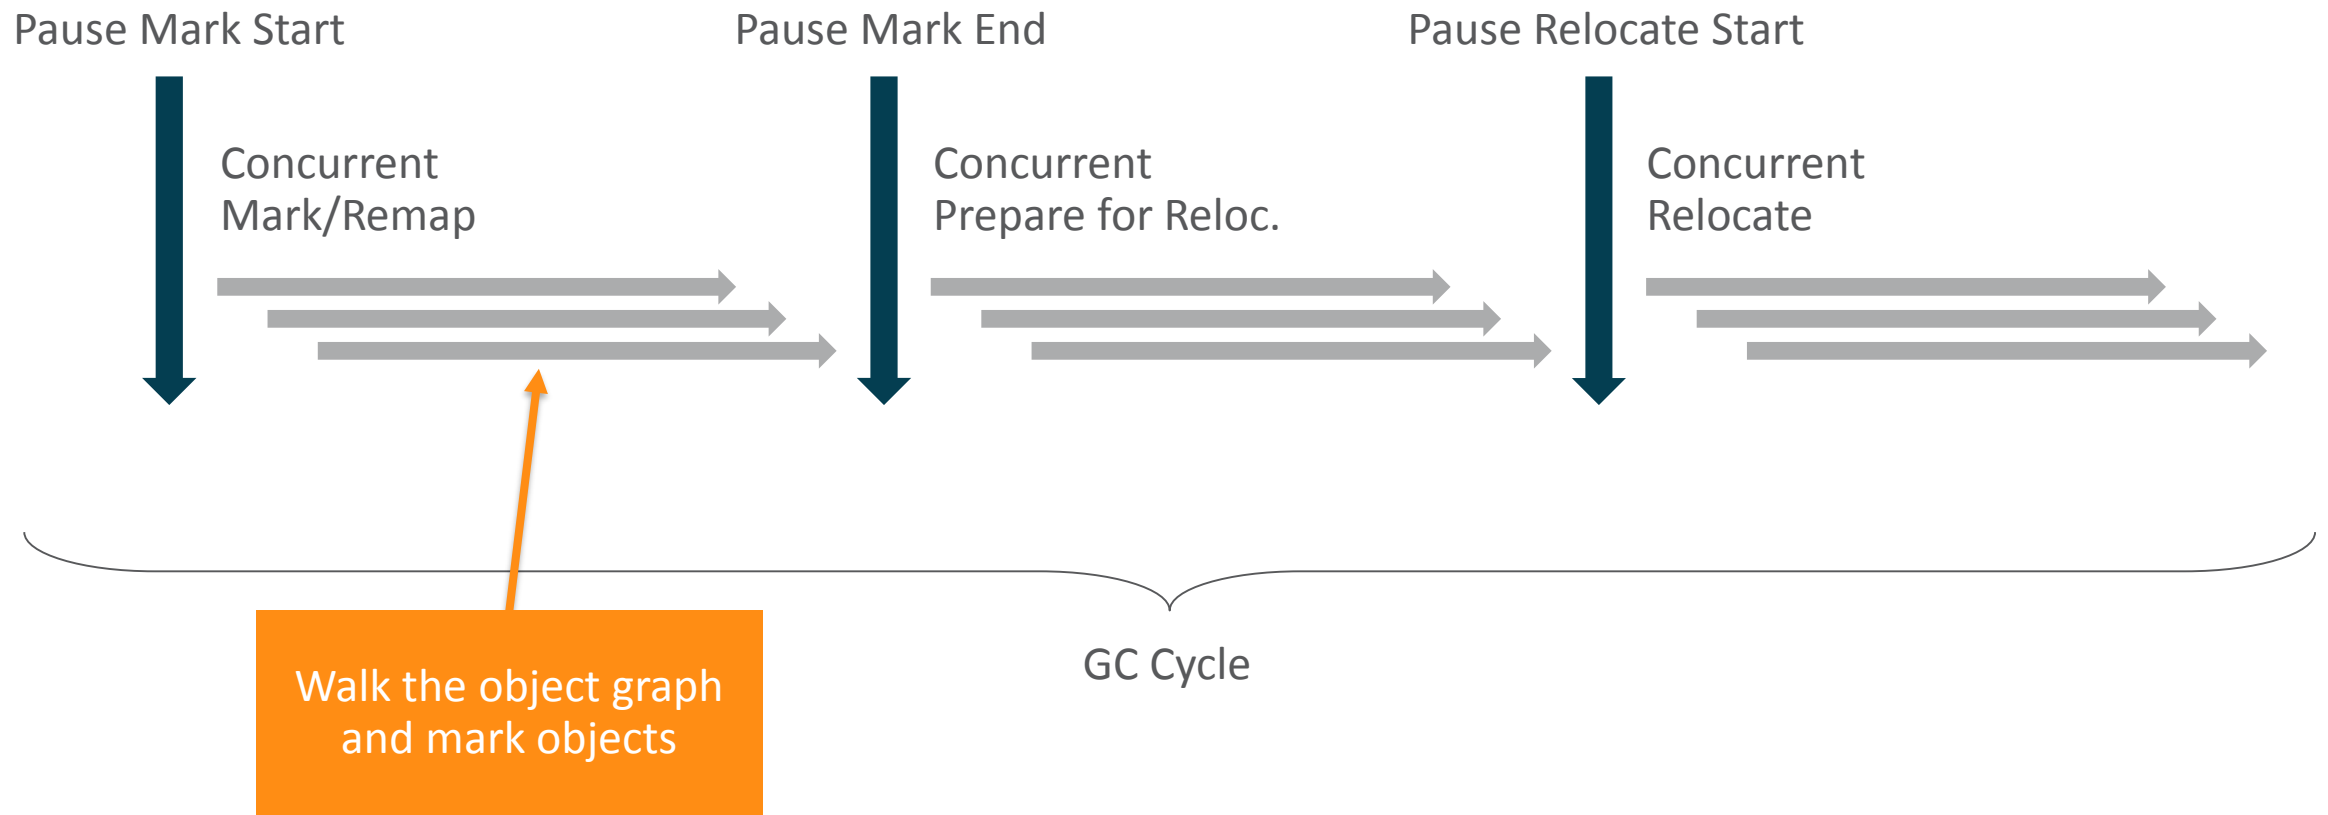
ZGC Phases



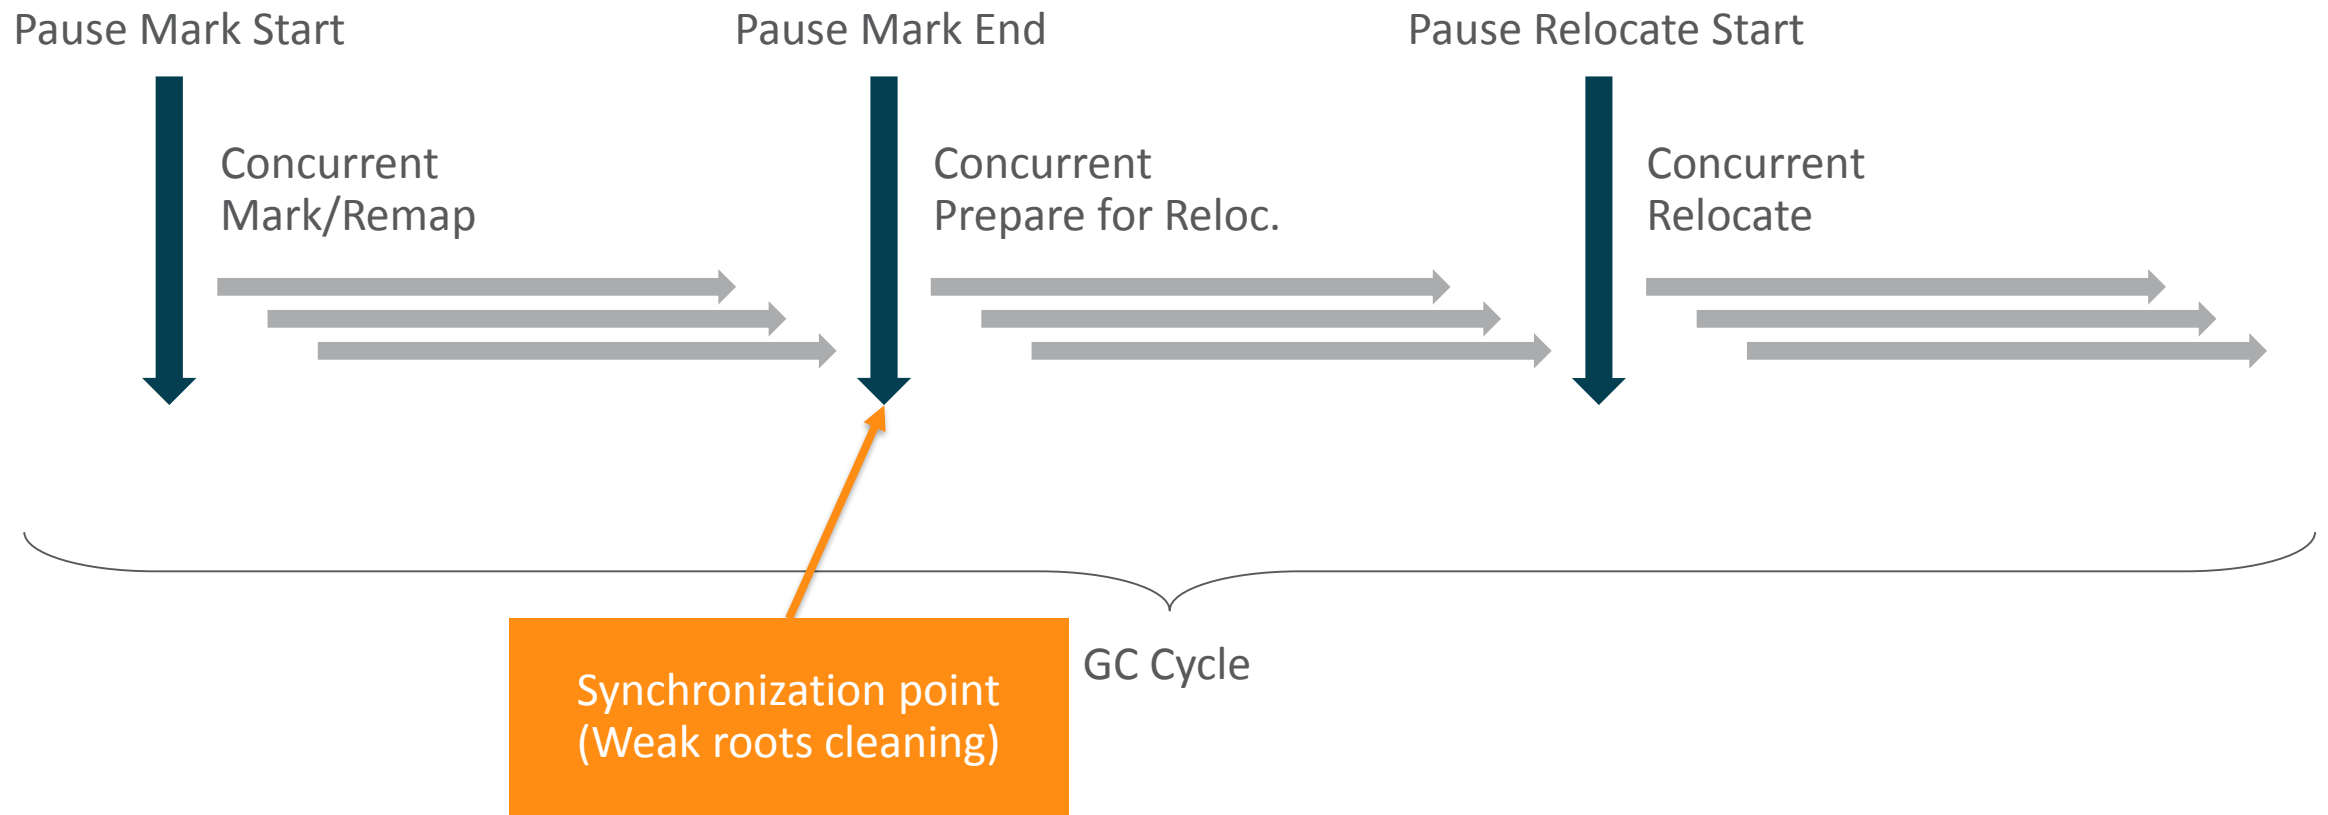
ZGC Phases



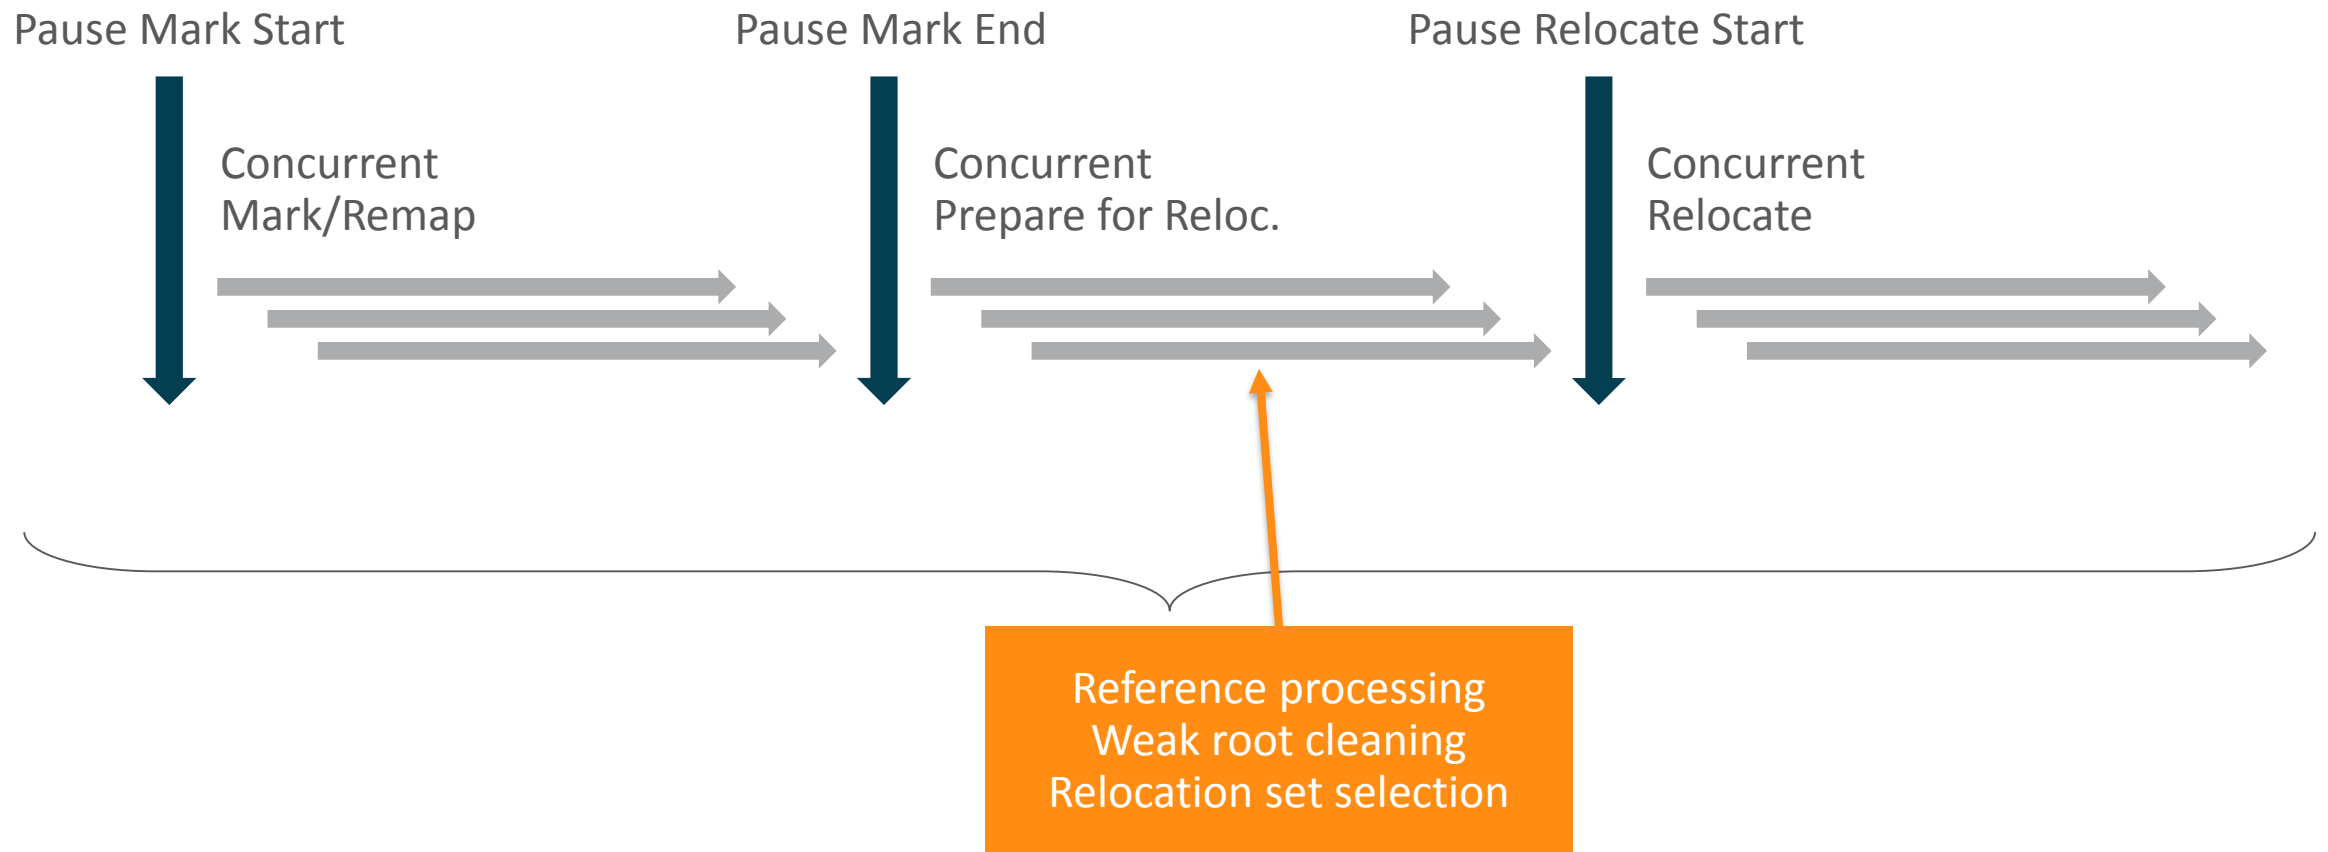
ZGC Phases



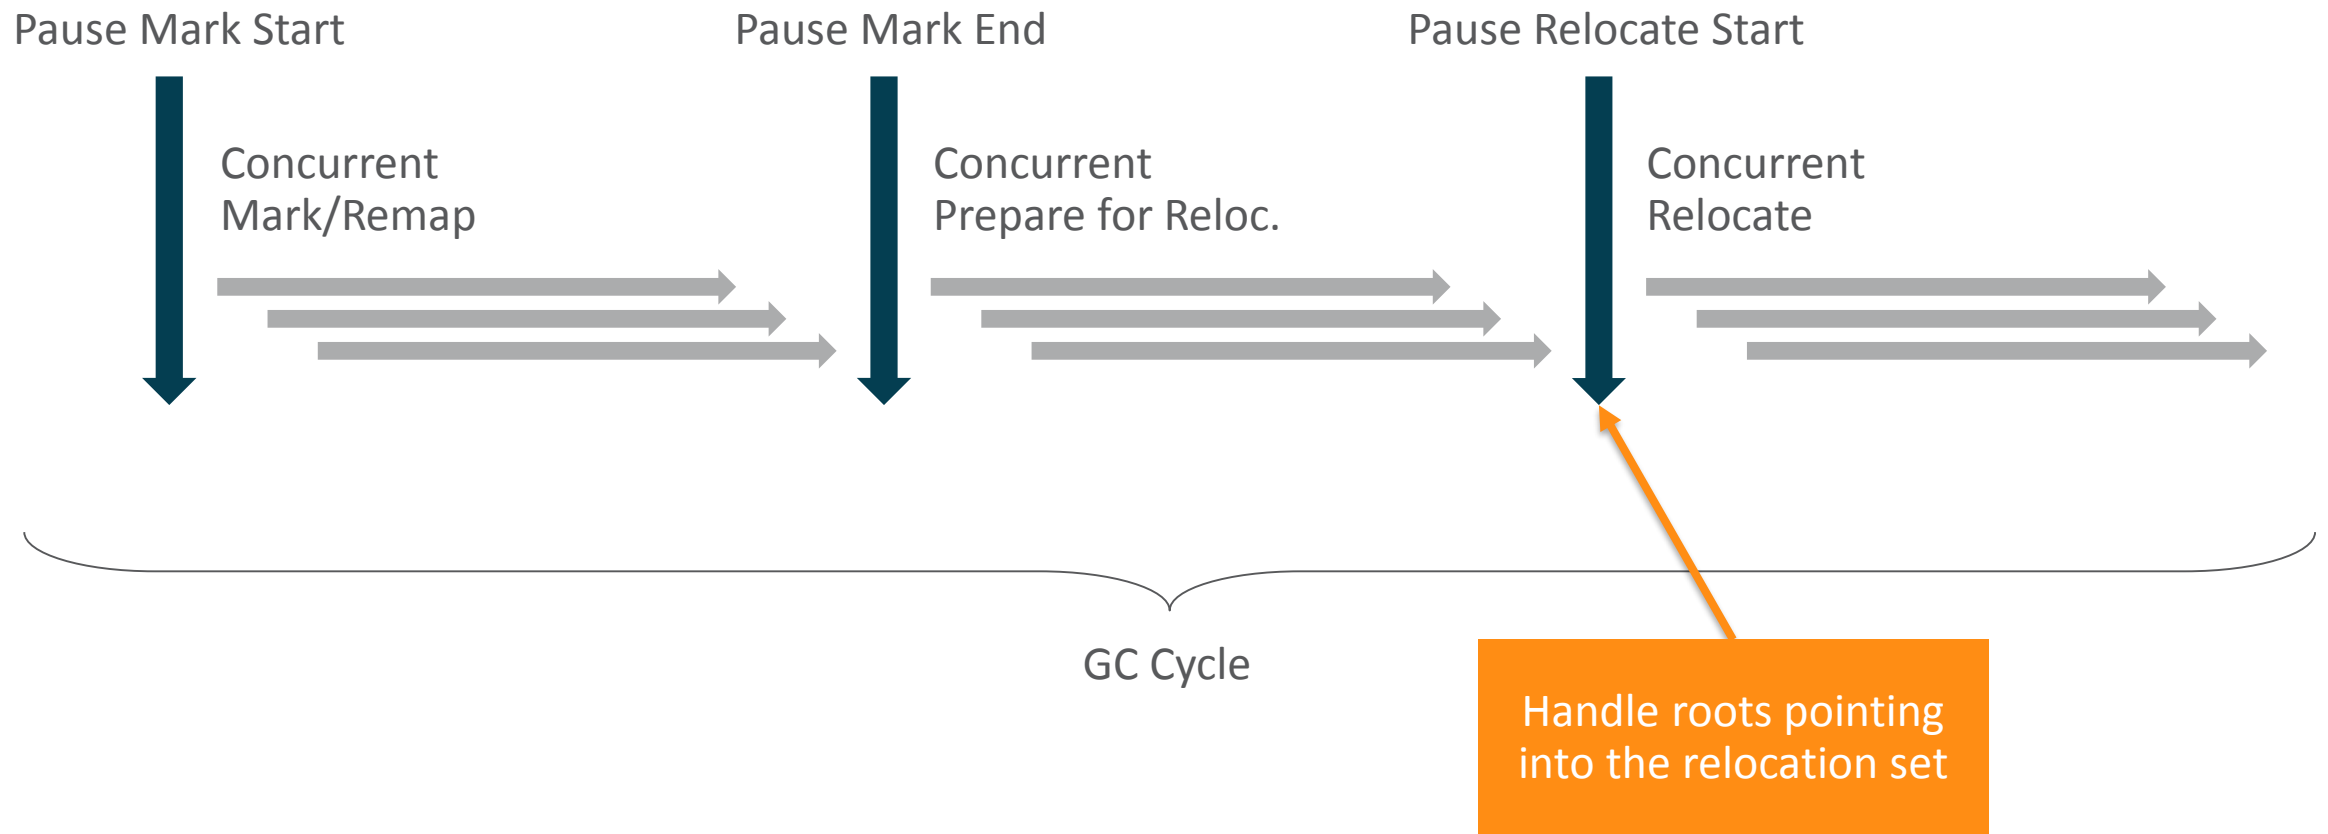
ZGC Phases



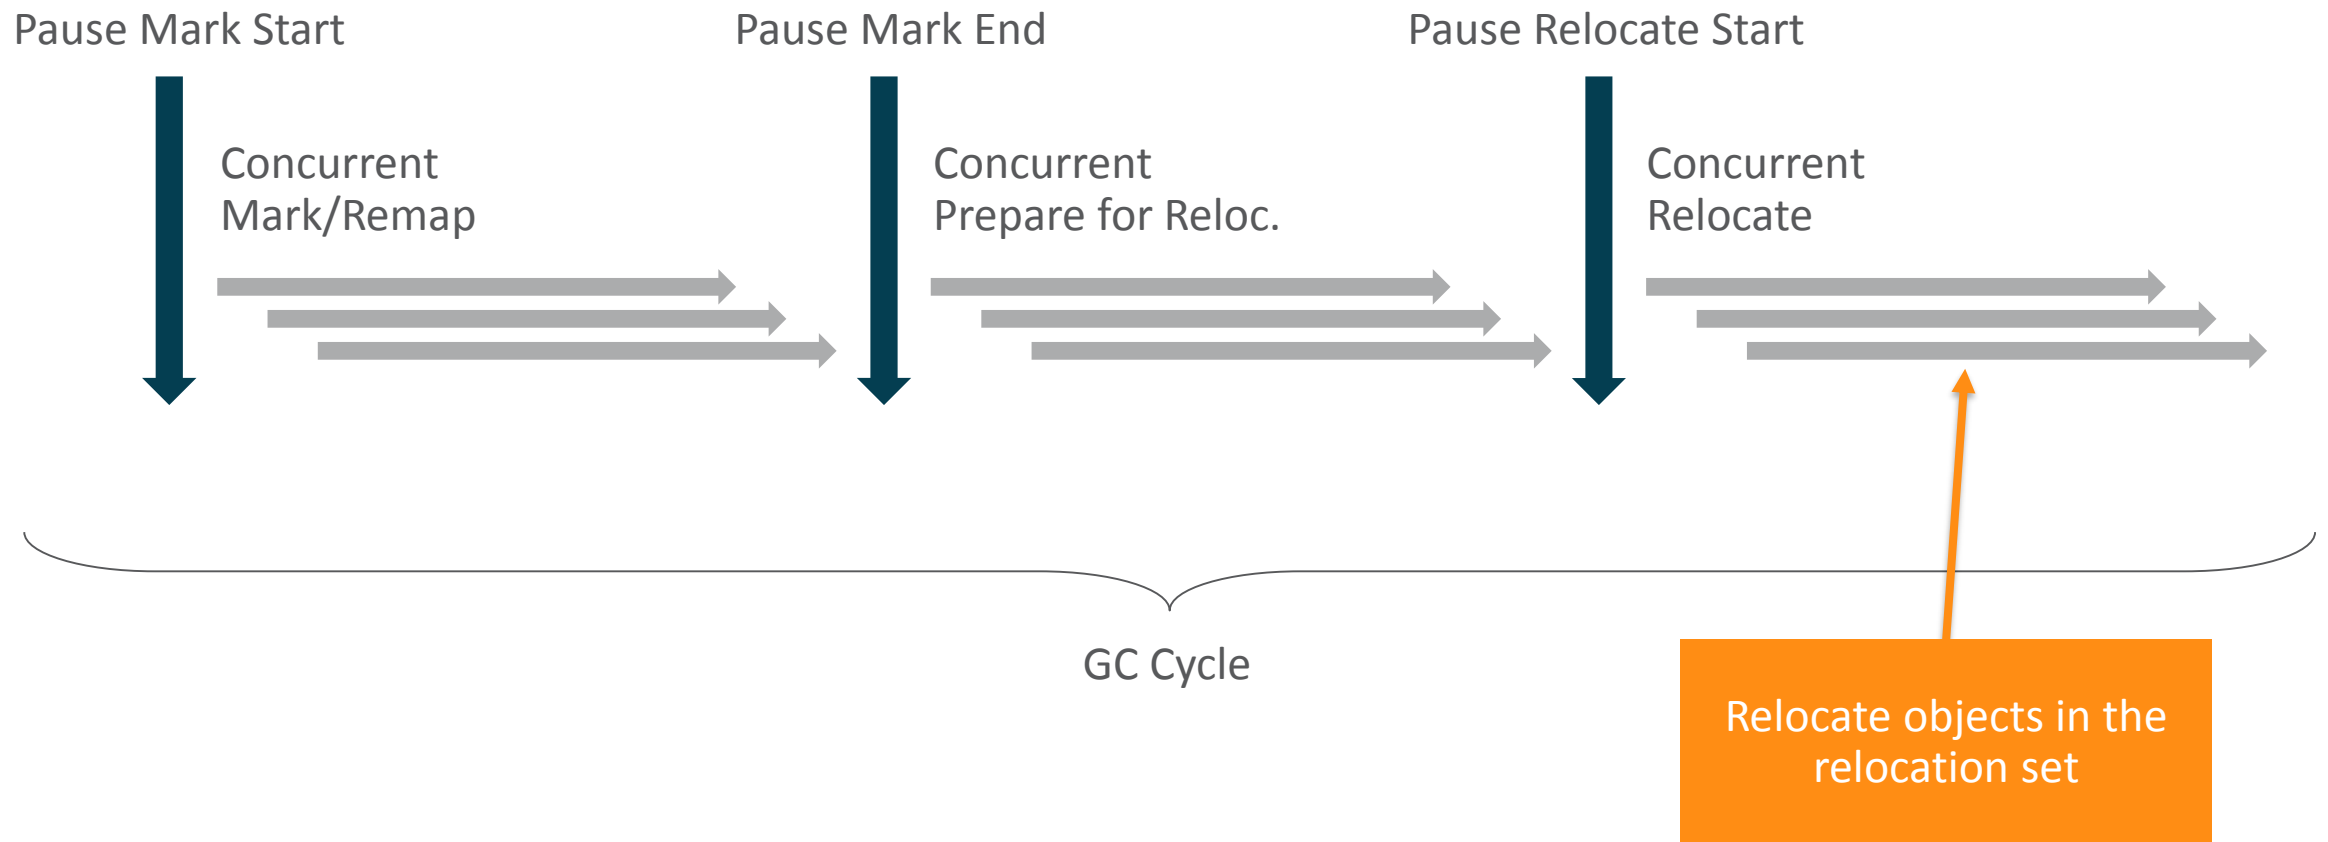
ZGC Phases



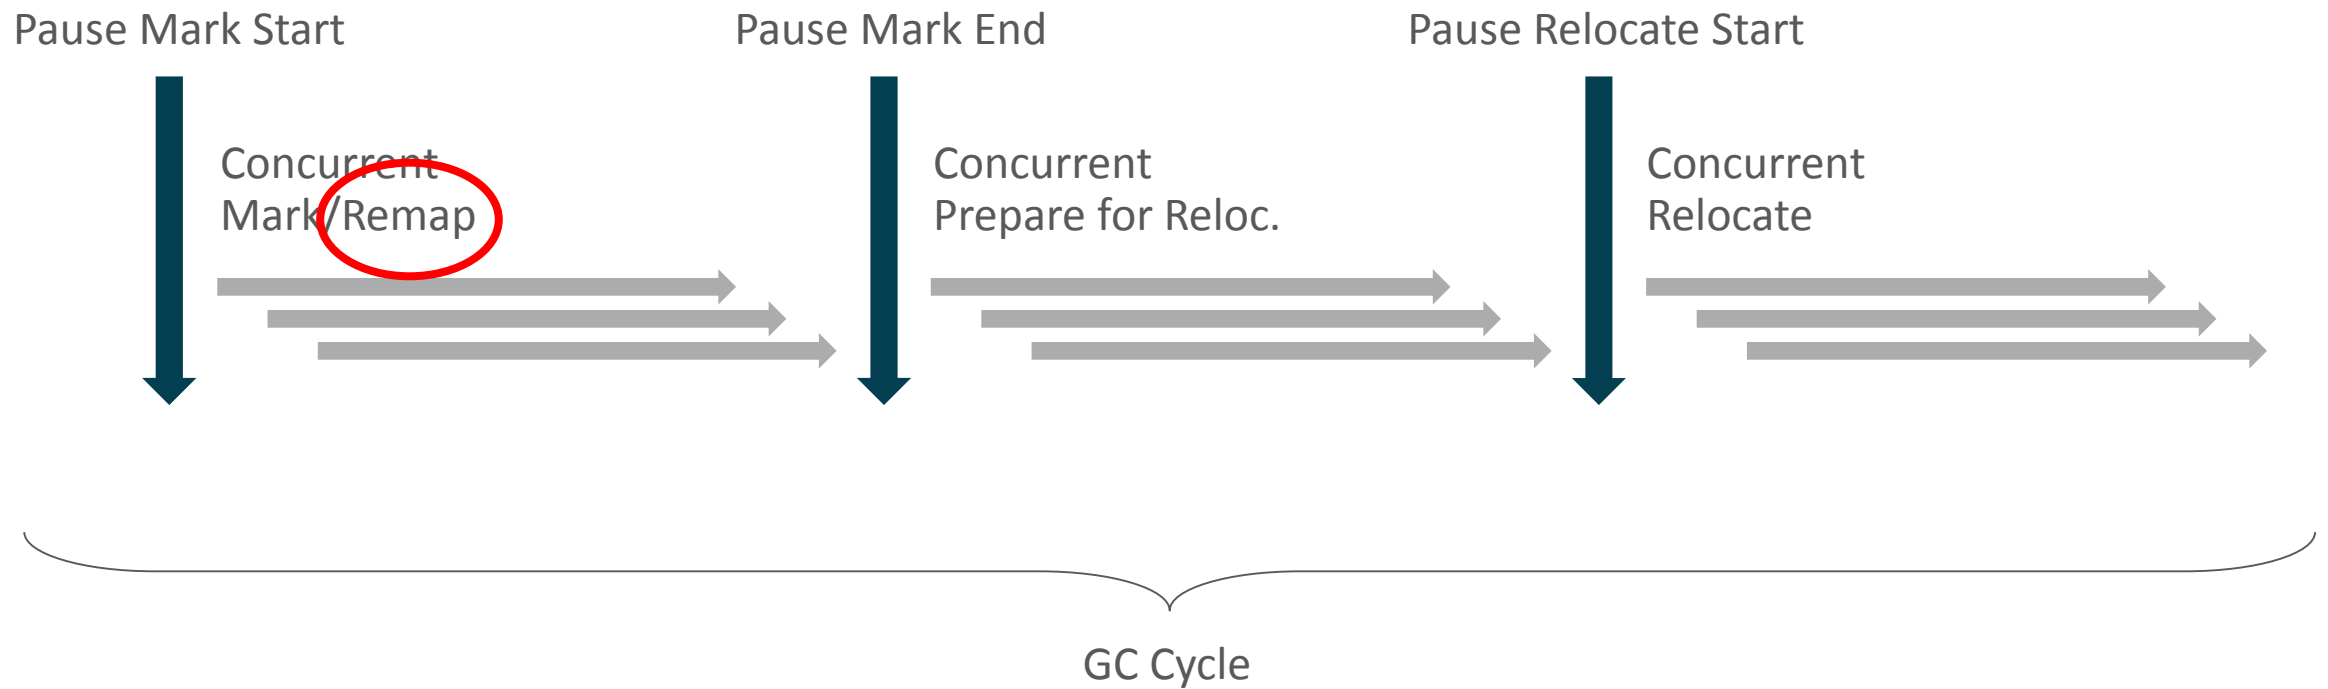
ZGC Phases



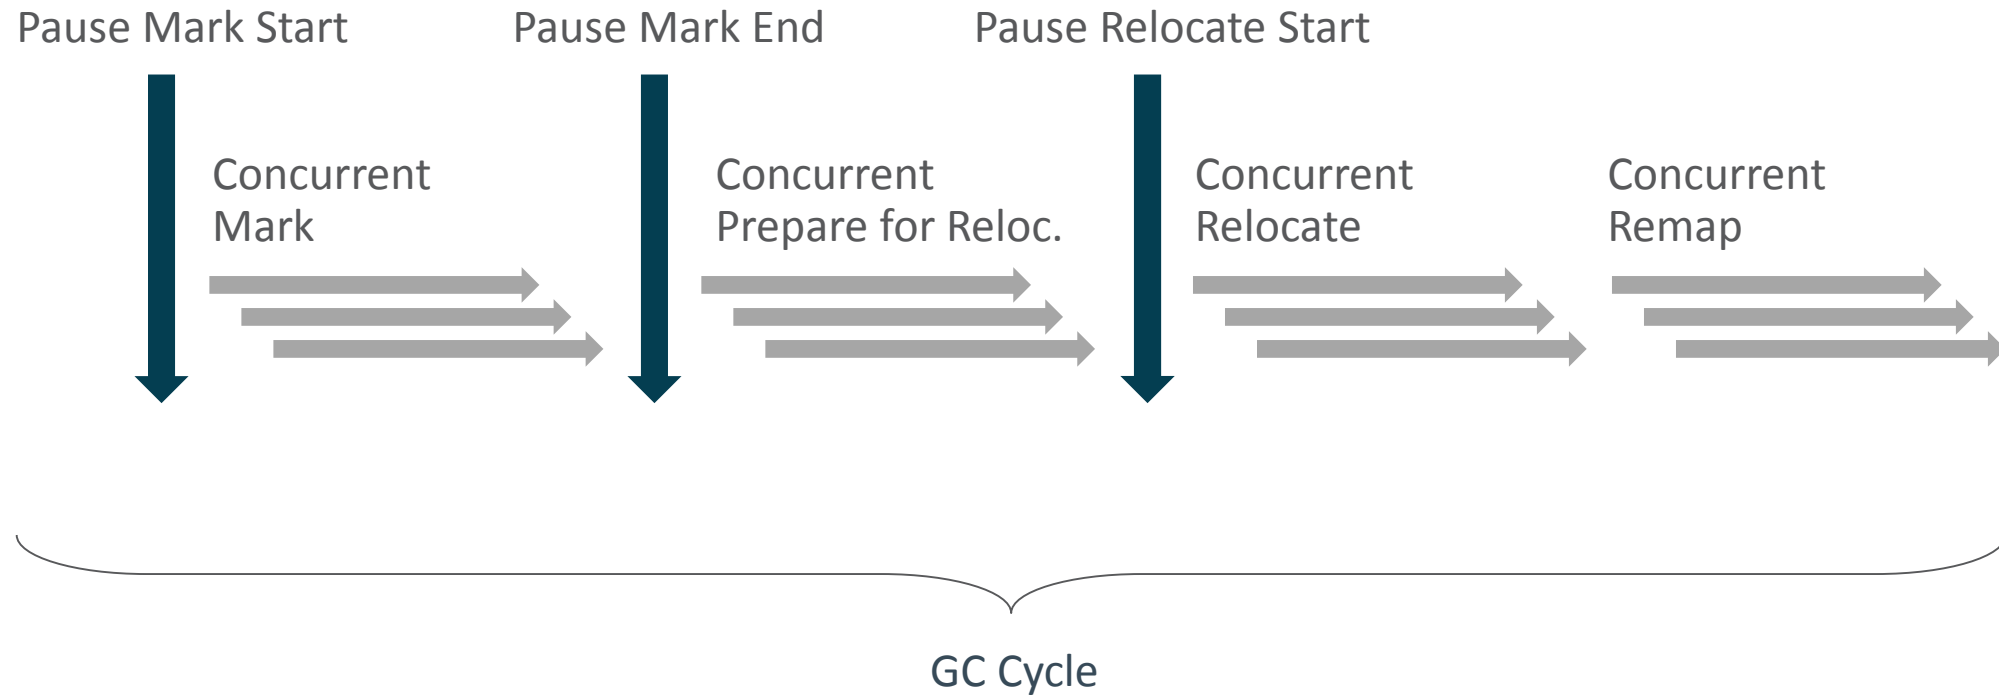
ZGC Phases



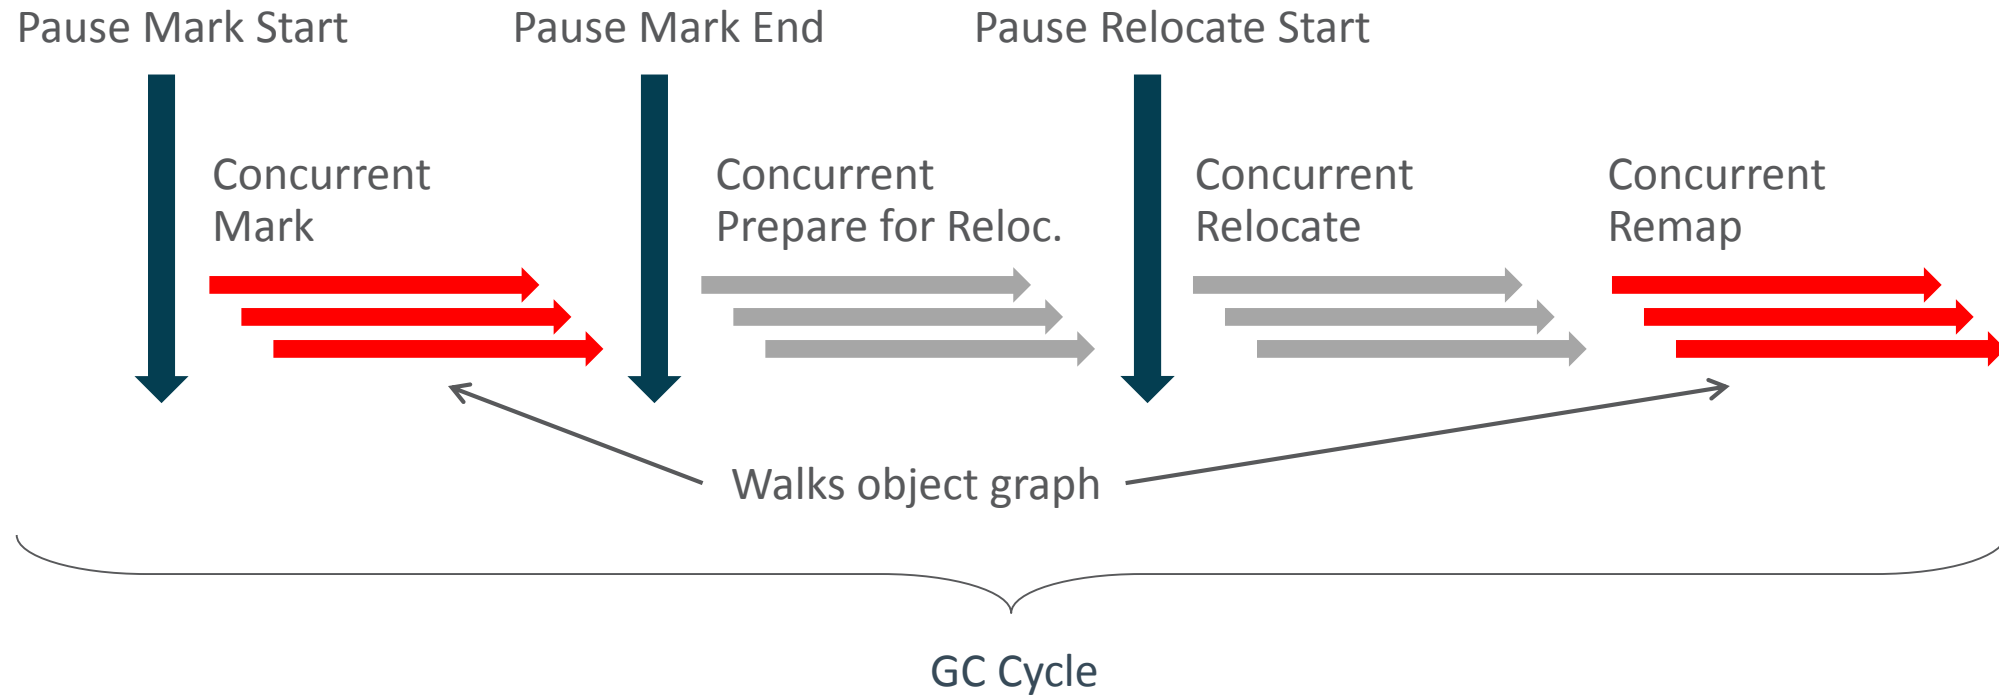
ZGC Phases




ZGC Phases




ZGC Phases




ZGC Phases



ZGC Phases



Heap Address Space



Heap Regions

Also known as ZPages

- Dynamically created/destroyed
- Dynamically sized
 - Multiple of **2MB** on **x86_64**
- Size groups
 - **Small** (2MB)
 - **Medium** (32MB)
 - **Large** (N x 2MB)




Colored Pointers



- Core design concept in ZGC
- **Metadata** stored in unused bits in 64-bit pointers
 - No support for 32-bit platforms
 - No support for CompressedOops
- **Virtual Address-masking** either in hardware, OS or software
 - Heap multi-mapping on Linux/x86_64
 - Supported in hardware on Solaris/SPARC


Colored Pointers

Layout on x86_64



Colored Pointers


Layout on x86_64



Colored Pointers

Layout on x86_64

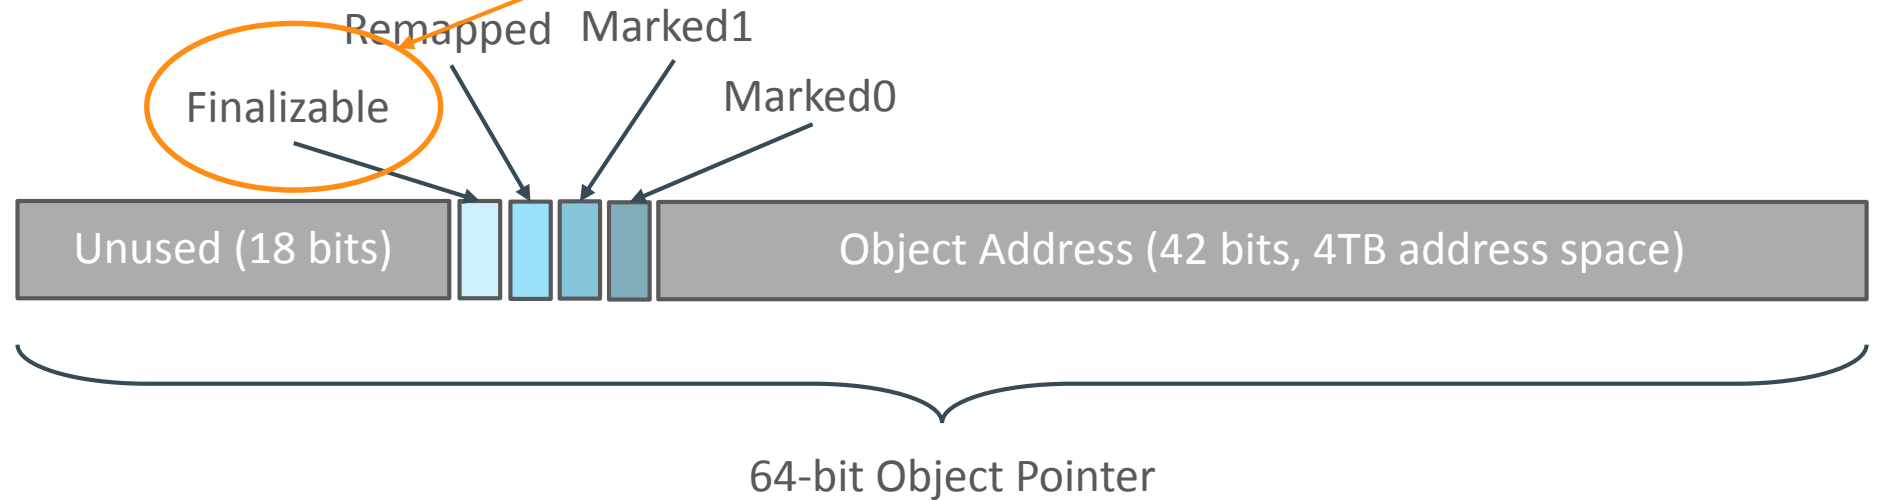
Known to not point into the relocation set?



Colored Pointers

Layout on x86_64

Only reachable through a Finalizer?



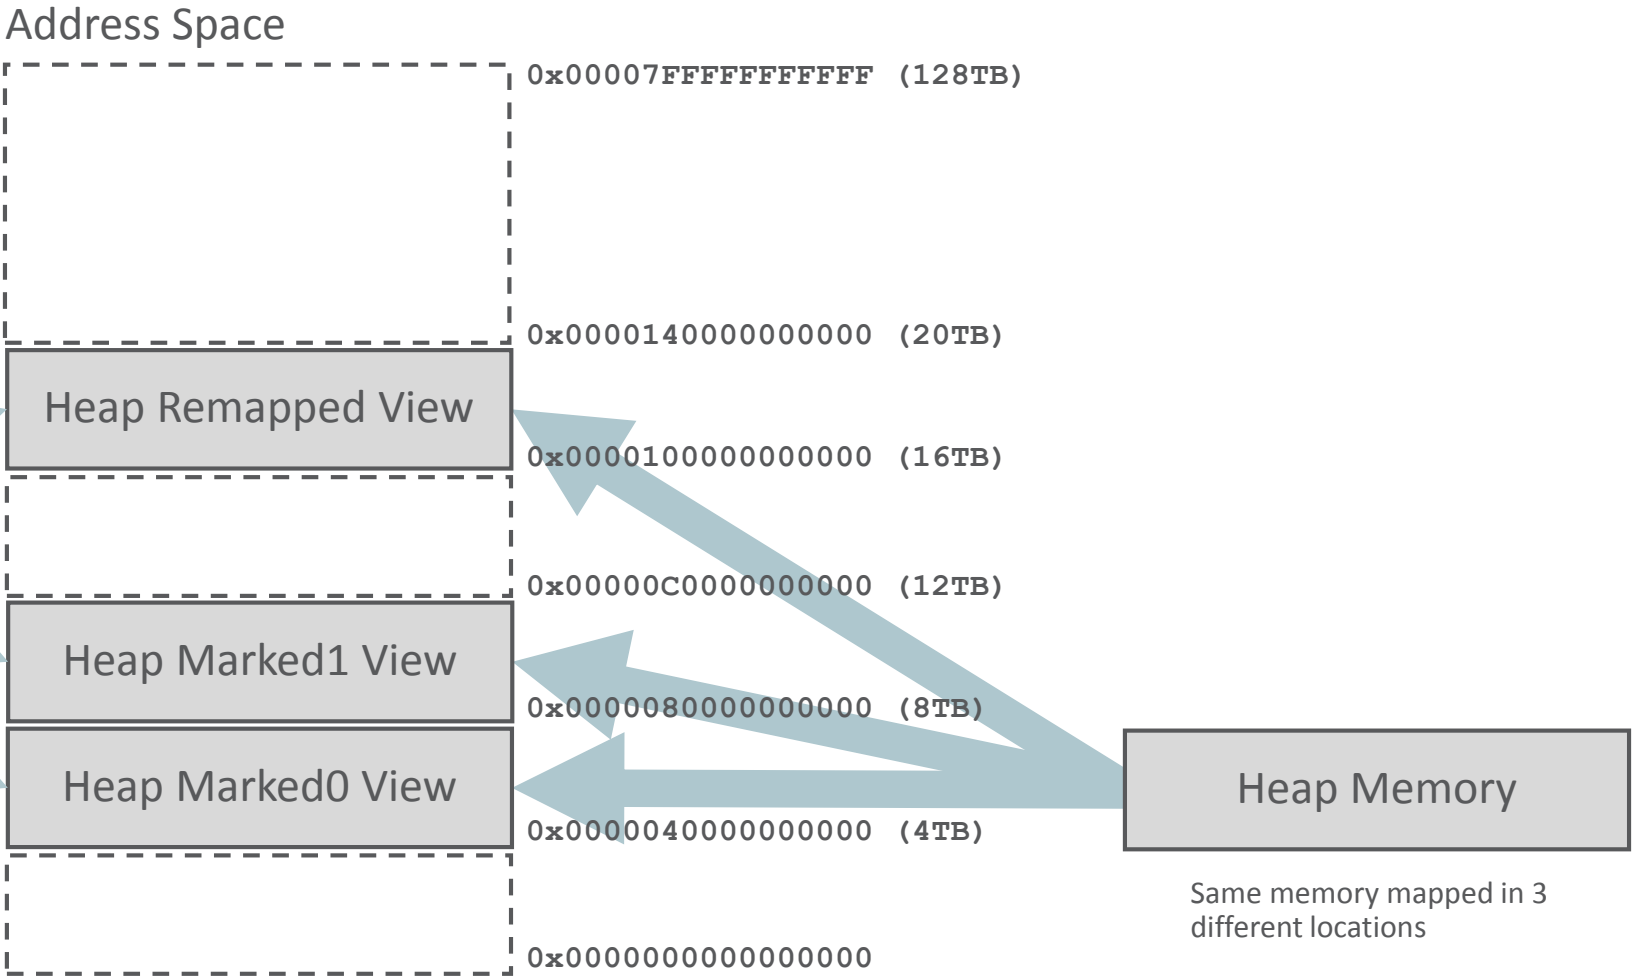
Heap Multi-Mapping on Linux/x86_64

Colorless pointer
0x00000000012345678

Colored pointer (Remapped)
0x00000100012345678

Colored pointer (Marked1)
0x00000800012345678


Colored pointer (Marked0)
0x00000400012345678



Heap Mapping on Solaris/SPARC

- Single heap mapping
- Virtual address masking in hardware
- **Load** and **store** instructions mask out metadata bits

Address Space



(ARM AArch64 also supports this)

Load Barrier

- Applied when **loading an object reference** from the heap
 - **Not** when later using that reference to access the object
 - Conceptually similar to the decoding of compressed oops
- Looks at the color of the pointer
 - Take action if the pointer has a “**bad**” color (mark/relocate/remap)
 - Change to the “**good**” color (repair/heal)
- Optimized for the common case
 - Most object references will have the “**good**” color

Load Barrier

```
Object o = obj.fieldA;           // Loading an object reference from heap
```

Load Barrier

```
Object o = obj.fieldA;  
<load barrier needed here>
```

```
// Loading an object reference from heap
```

Load Barrier

```
Object o = obj.fieldA;  
<load barrier needed here>  
Object p = o;  
o.doSomething();  
int i = obj.fieldB;
```

```
// Loading an object reference from heap  
  
// No barrier, not a load from heap  
// No barrier, not a load from heap  
// No barrier, not an object reference
```

Load Barrier

```
Object o = obj.fieldA;  
<load barrier needed here>
```

```
// Loading an object reference from heap
```

Load Barrier

```
Object o = obj.fieldA;           // Loading an object reference from heap  
load_barrier(register_for(o), address_of(obj.fieldA));
```


Load Barrier

```
Object o = obj.fieldA;           // Loading an object reference from heap
if (!(o & good_bit_mask)) {
    if (o != null) {
        slow_path(register_for(o), address_of(obj.fieldA));
    }
}
```

Load Barrier

```
Object o = obj.fieldA;           // Loading an object reference from heap
if (o & bad_bit_mask) {
    slow_path(register_for(o), address_of(obj.fieldA));
}
```

Load Barrier

```
mov    0x20(%rax), %rbx
test   %rbx, (0x16)%r15
jnz    slow_path
```

```
// Object o = obj.fieldA;
// Bad color?
// Yes -> Enter slow path and
// mark/relocate/remap, adjust
// 0x20(%rax) and %rbx
```

Load Barrier

```
mov    0x20(%rax), %rbx
test   %rbx, (0x16)%r15
jnz    slow_path
```

```
// Object o = obj.fieldA;
// Bad color?
// Yes -> Enter slow path and
// mark/relocate/remap, adjust
// 0x20(%rax) and %rbx
```

~4% execution overhead on **SPECjbb[®]2015**

Load Barrier (r12 version)

```
mov    0x20(%rax), %rbx
test   %rbx, %r12
jnz    slow_path
```


```
// Object o = obj.fieldA;
// Bad color?
// Yes -> Enter slow path and
// mark/relocate/remap, adjust
// 0x20(%rax) and %rbx
```

Always keep **bad_bit_mask** in r12

- Avoids a memory load, but reserves a register
- We don't support compressed oops, so we can repurpose r12, the heap base register

Mark

- Concurrent & Parallel
- Load barrier
 - Detects loads of non-marked object pointers
- Finalizable mark
 - Enabler for Concurrent Reference Processing
- Thread local handshakes
 - Used to synchronize end of concurrent mark
- Striped



Striped Mark



- Scalability
 - Heap divided into logical stripes
 - Isolate each GC thread to work on its own stripe
 - Minimized shared state
- **Edge** pushing vs. **Node** pushing
 - Potentially more work
 - ... but lends itself better to parallel processing

Striped Mark



Heap


GC
Thread 0

GC
Thread 1

GC
Thread 2

GC
Thread 3

Striped Mark



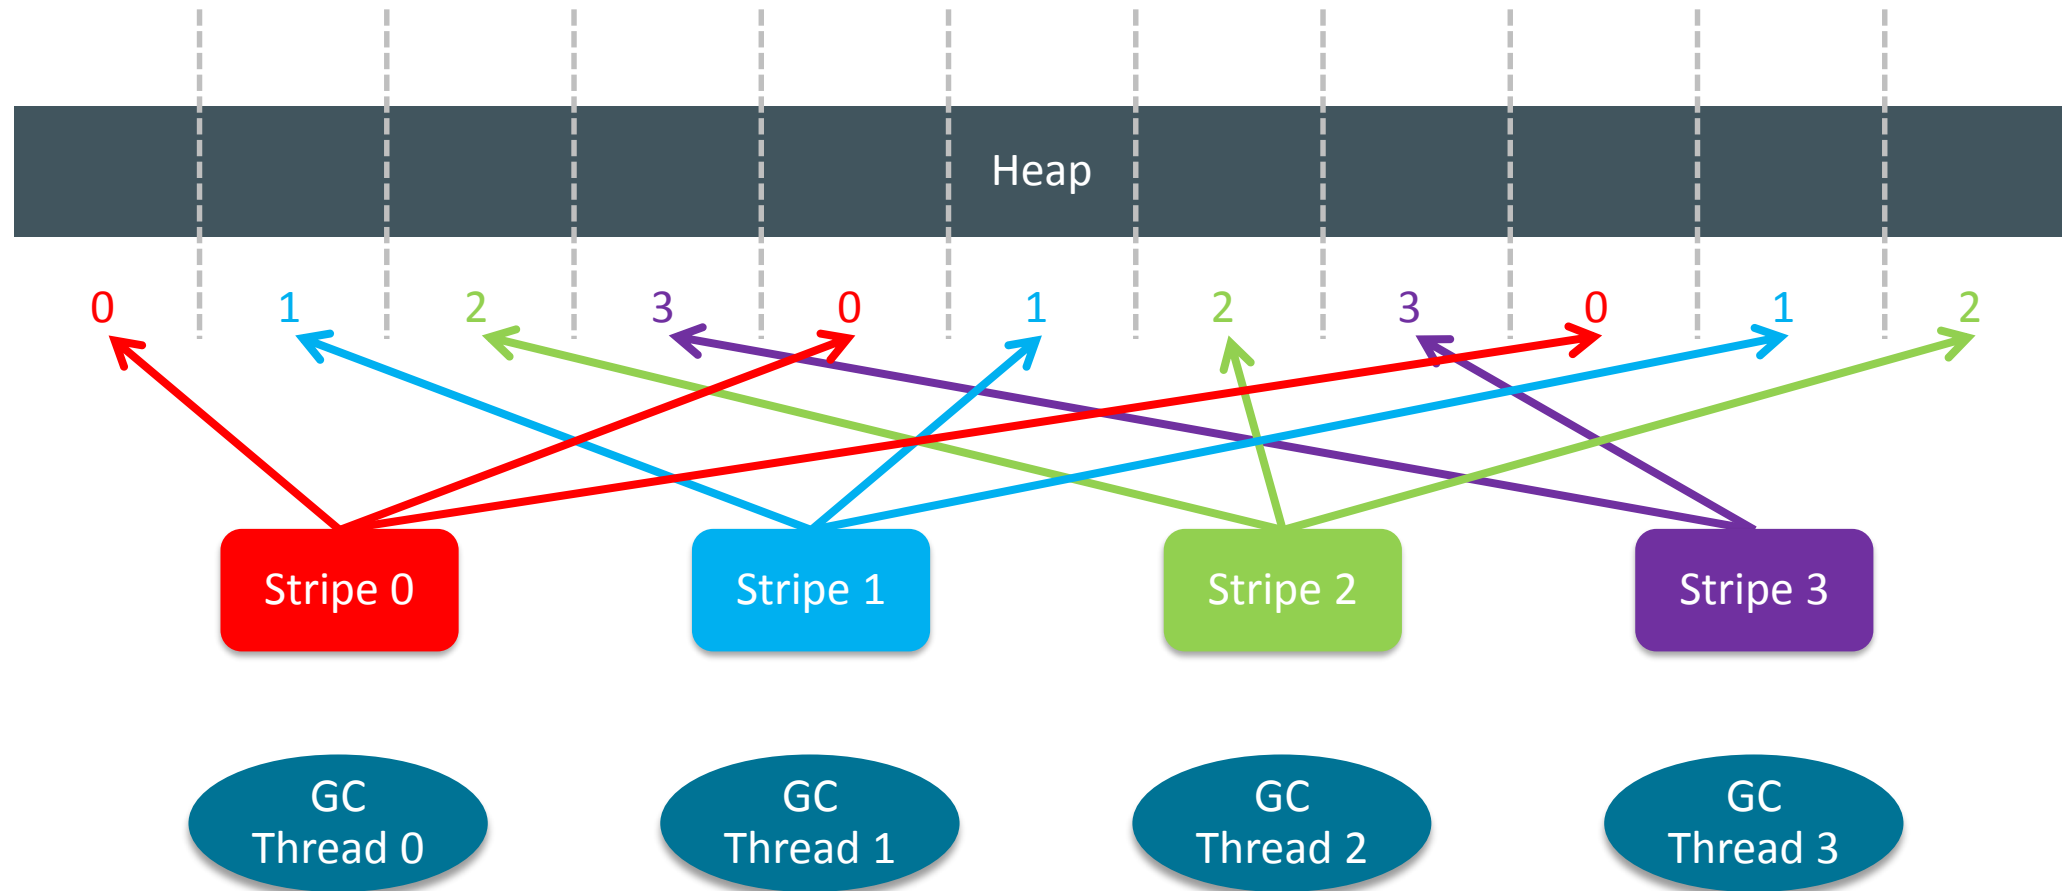
GC
Thread 0

GC
Thread 1

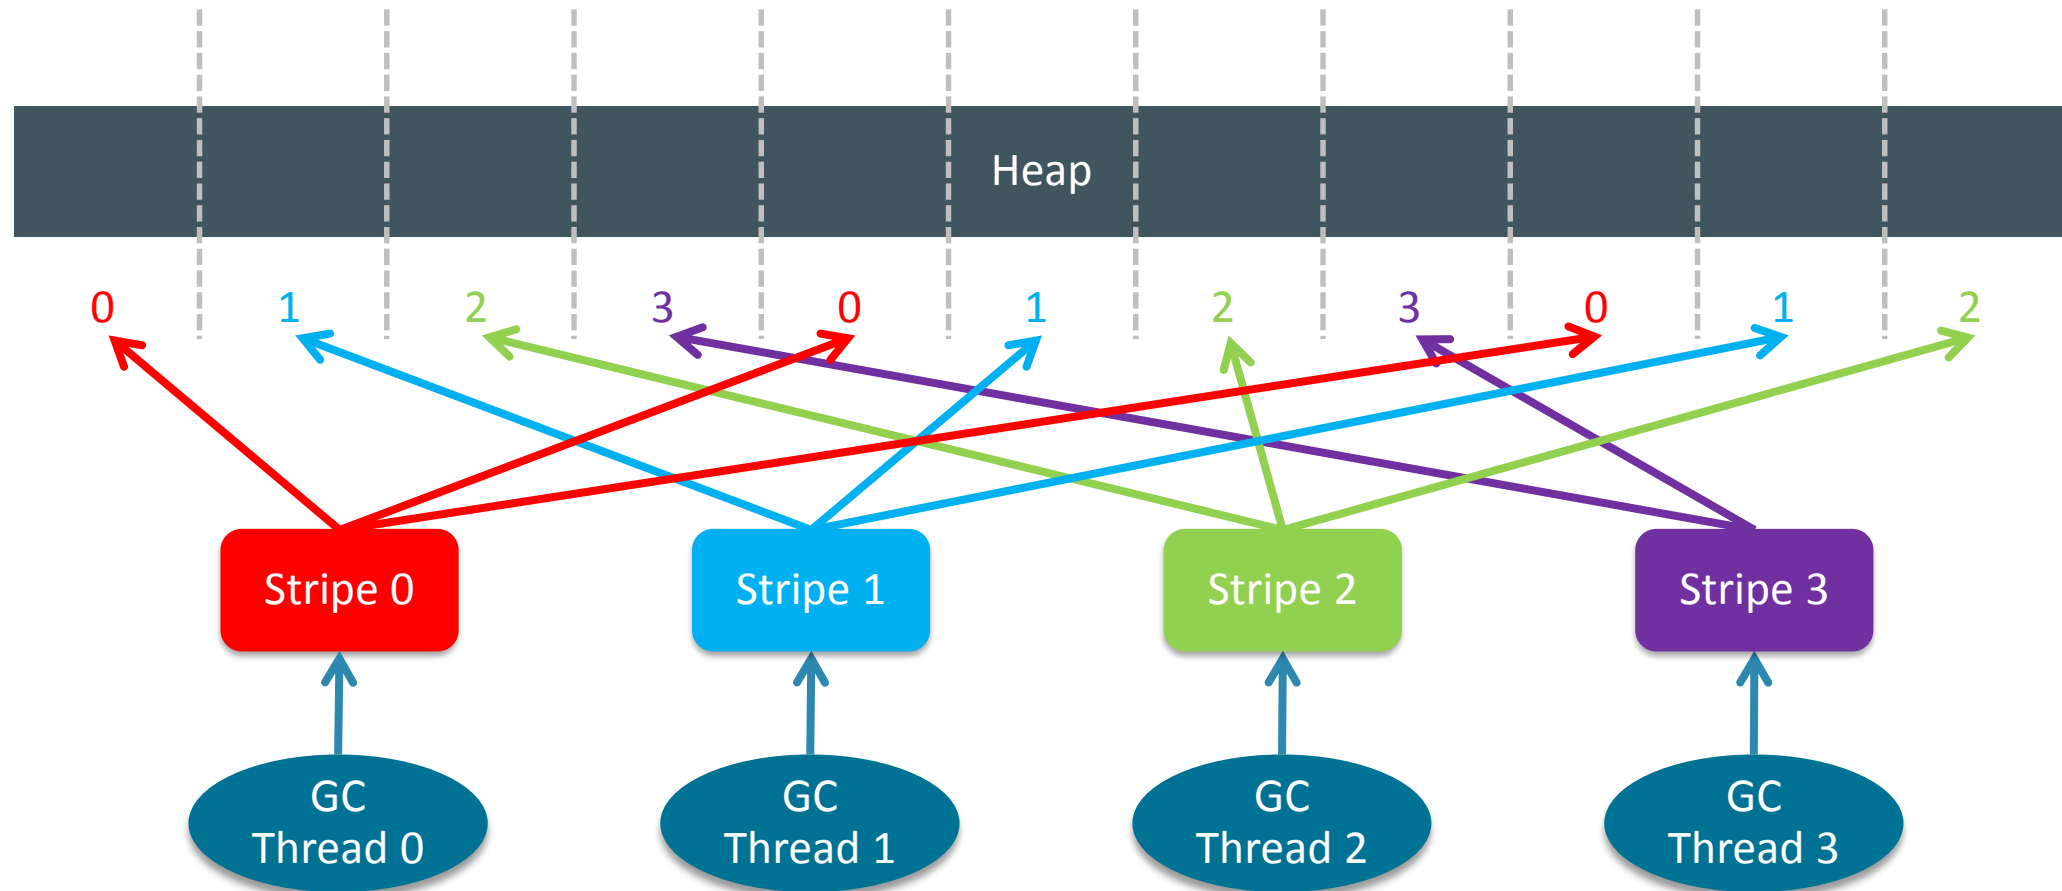
GC
Thread 2

GC
Thread 3

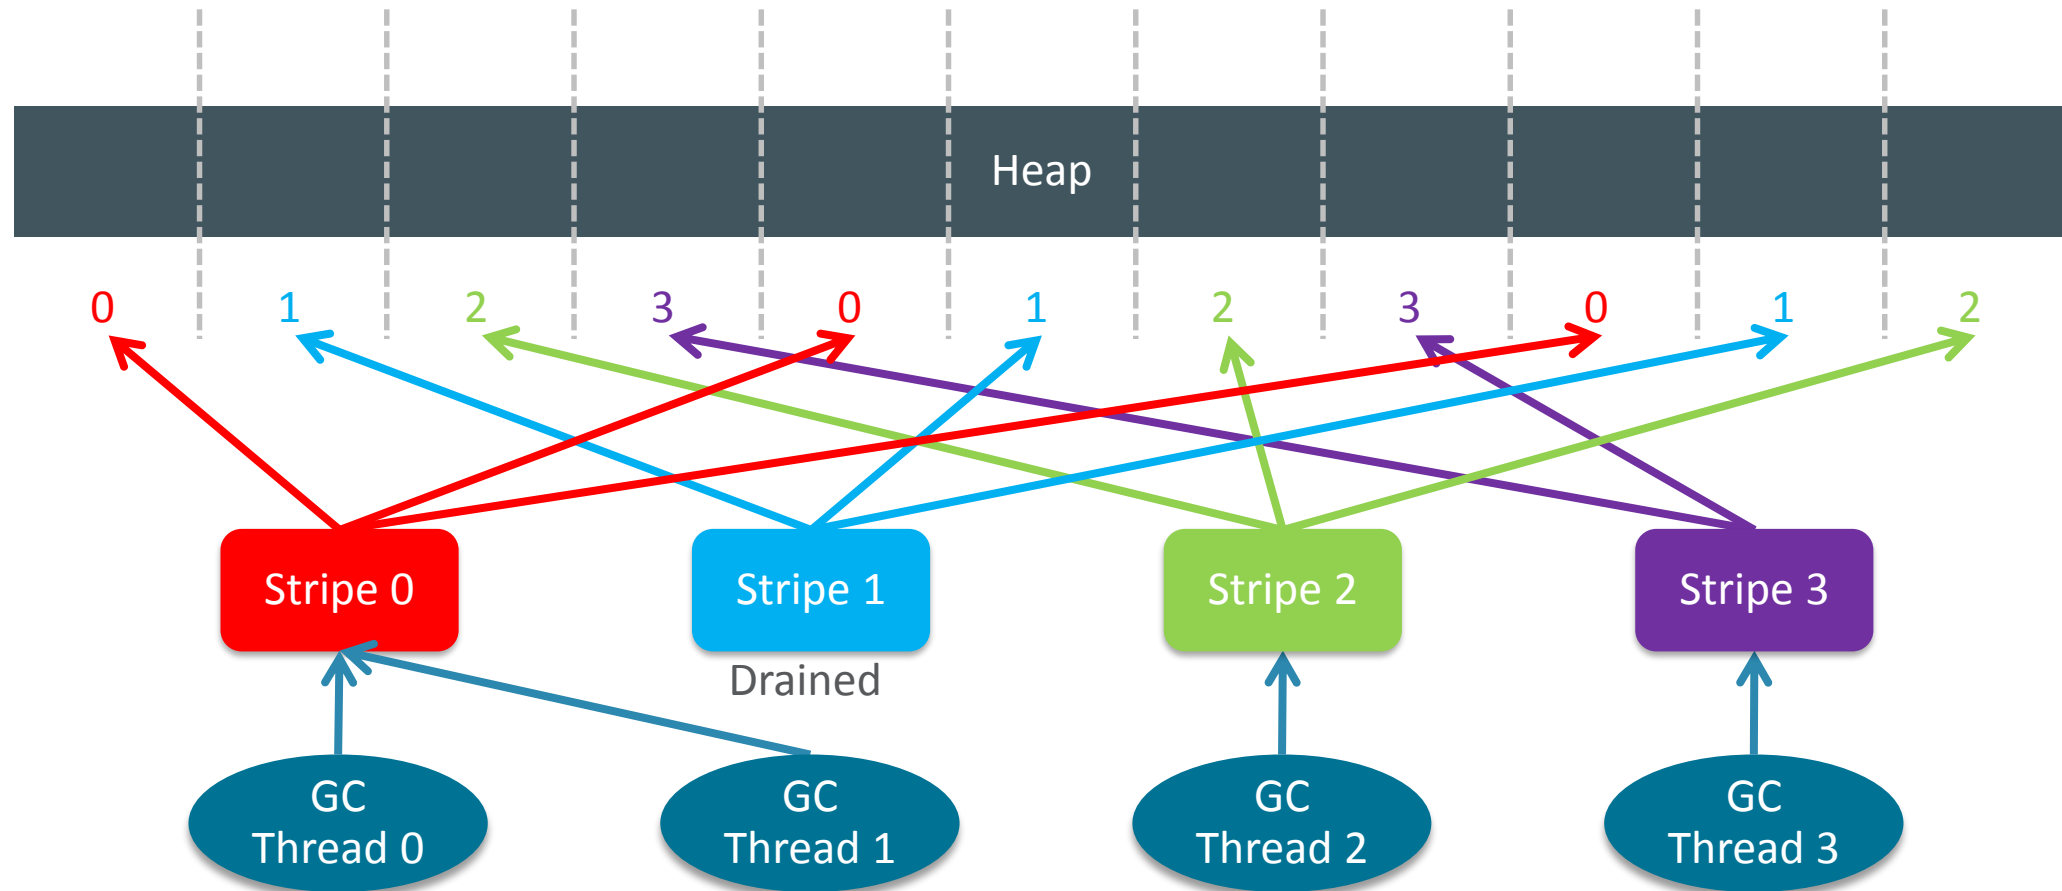
Striped Mark



Striped Mark



Striped Mark




Reference Processing

Dealing with Soft/Weak/Final/PhantomReference


- Concurrent & Parallel
- Liveness/Reachability analysis
 - **Complete** after concurrent mark
 - Strongly reachable, Final reachable and Unreachable
- Processing/Enqueuing
 - **Single** pass
 - Load barrier **blocks** resurrection attempts (e.g. through **Reference.get()**)

Relocation




- Concurrent & Parallel
- Load barrier
 - Detects loads of object pointers pointing into the relocation set
 - Java threads help out with relocation if needed
- Off-heap forwarding tables
 - No forwarding information stored in old copies of objects
 - Important for immediate reuse of heap memory

GC Cycle Example




Pause Mark Start

 Marked




Pause Mark Start

 Marked




Pause Mark Start

 Marked




Pause Mark Start

 Marked




Concurrent Mark

 Marked




Concurrent Mark

 Marked




Concurrent Mark

 Marked




Concurrent Mark

 Marked




Pause Mark End

 Marked




Concurrent Prepare for Relocate

 Marked





Concurrent Prepare for Relocate



 Marked




Pause Relocate Start



Pause Relocate Start




Pause Relocate Start





Pause Relocate Start

Marked



Remapped + Relocated





Concurrent Relocate





Concurrent Relocate





Concurrent Relocate




Concurrent Relocate




Concurrent Relocate




GC Cycle Completed




GC Cycle Completed




GC Cycle Completed




Pause Mark Start (Second Cycle)




Pause Mark Start (Second Cycle)




Pause Mark Start (Second Cycle)




Pause Mark Start (Second Cycle)





Concurrent Mark (Second Cycle)




Concurrent Mark (Second Cycle)




Concurrent Mark (Second Cycle)





Concurrent Mark (Second Cycle)



Pause Mark End (Second Cycle)



Concurrent Prepare for Relocate (Second Cycle)



Forwarding Tables Freed

Agenda

- 1 ➤ What is ZGC?
- 2 ➤ Some Numbers
- 3 ➤ Under The Hood
- 4 ➤ Going Forward**
- 5 ➤ How To Get Started

In The Works

- GC Barrier API
 - Make it **easier** to plug in new GCs (ZGC, Shenandoah, Epsilon)
- Concurrent class unloading & weak roots
 - Traditionally done in a Stop-The-World pause
 - Impacts **JITs** and **Runtime** subsystems
- Addressing non-GC induced latencies
 - Time to safepoint/unsafepoint, object monitor deflation, etc.



Foundation for Future GC Features

Colored Pointers + Load Barriers

- Thread local GC scheme
- Track heap access patterns
- Use non-volatile memory for rarely used parts of the heap
- Compress or archive parts of the heap
- Object properties encoded in pointers
- Allocation tricks
- etc.



Agenda

- 1 ➤ What is ZGC?
- 2 ➤ Some Numbers
- 3 ➤ Under The Hood
- 4 ➤ Going Forward
- 5 ➤ How To Get Started**

How To Get Started

Download

- Official **early access builds** will be available soon-ish, but until then...
- Download & build

```
$ hg clone http://hg.openjdk.java.net/zgc/zgc
$ cd zgc
$ sh configure
$ make images
```

- Run

```
$ ./build/linux-x86_64-<...>/images/jdk/bin/java
```

How To Get Started

JVM Options

- Enable ZGC: **-XX:+UseZGC**
- Tuning
 - If you care about latency, do **not** overprovision your machine
 - Max heap size: **-Xmx<size>**
 - Number of concurrent GC threads: **-XX:ConcGCThreads=<number>**
- Logging
 - Basic logging: **-Xlog:gc**
 - Detailed logging useful when tuning: **-Xlog:gc***

Feedback Welcome!

<http://wiki.openjdk.java.net/display/zgc/>

OpenJDK

